

---

# X86 ASSEMBLY

---

A PREPRINT

**Angel Aguayo**  
University of Arizona  
aaguayo30@email.arizona.edu

**Nicholas Kunzler**  
University of Arizona  
nkkunzler@email.arizona.edu

April 29, 2020

## ABSTRACT

x86 assembly language is a verbose assembly language used in most servers and workstation architectures today. x86 follows a CISC instruction architecture which allows for complex and powerful instructions not available on RISC architectures. Although the use of programming in an assembly language, such as x86, is extremely limited due to the large time commitment required, it is an excellent language for extreme performance optimizations and lower level applications all due to how close assembly matches actual machine instructions. Although x86 and other assembly languages consist of very simple instruction sets and may seem limited, it is capable of creating complex programs, which can be seen by all the programs compiled to x86.

## 1 Introduction

x86 is a CISC, Complex Instruction Set Computing, assembly language with a variety of syntax styles and architectures. x86 conforms to 16-bit, 32-bit, and 64-bit architectures with the backwards compatibility with smaller bit architectures. Unlike RISC, Reduced Instruction Set Computing, x86 utilizes complex instructions, instructions that utilize multiple other simplified instructions, to create powerful new instructions. Although CISC architectures provided these powerful new instructions, this is not always desired. RISC instruction sets, such as ARM, tend to consume less power and therefore are used more in phones or smaller devices as the power draw should be reduced. x86 is used in most servers and workstations as these applications do not have the concern of running on battery or limited power sources that would be overwhelmed with complex instructions [3]. Overall, x86 provides eliminates the luxuries of more modern programming languages in order to reach a level as close to the machine as possible while also having the ability to create arbitrarily complex programs.

## 2 History

Intel was founded by two American engineers in 1968 with 2.5 million dollars in funding. Before they developed the x86 architecture, they focused on developing DRAM memory chips. After a series of unsuccessful and successful chips and becoming a public company, Intel pivoted towards microprocessors, beginning with the 4004 in 1971, which was primarily used within Japanese calculators. Following up in 1972, Intel released the 8 bit CPU called the 8008, and eventually the 8080 2 years later. These both used a simplistic instruction set architecture (ISA), with 8 registers and 8-bit instructions. In 1978, Intel altered its ISA with the 8086 microprocessor, using 80x86, known commonly as x86. With 16-bit instructions and 1 megabyte of main memory, this is faster than any of the previous microprocessors. Development began in it in 1976, and it was originally meant as a side project as they were planning to jump to 32-bit with the 8800 they were working on. The technology at the time inhibited the 8800 from continuing, so they hired Stephen Morse, a software and electrical engineer who identified key design flaws with the 8800, to be the sole designer of the 8086. This was the first time they have hired a software engineer to design their microprocessors. With a new software oriented approach of what features to add to make the software more efficient instead of what features can be added in, Intel ended up revolutionizing the industry [1].

At launch, the 8086 gained only minor traction due to the Z80, their competitors chip, being in most business machines as the standard. Eventually it entered the market of embedded applications as NASA used it for controlling their

diagnostic tests. After Morse left Intel in 1979, Intel released the 8088, which drew the attention of IBM, who chose the use the 8088 for their first mass produced personal computer (PC). The 8088 was backwards compatible, where it would send 16 bits out in 8 bit cycles, allowing the x86 ISA to be continued. IBM's PC used off the shelf parts, and as Intel's 8088 was 16-bit and could reduce chip count, IBM went with them. Over time, other companies started cloning IBM's PC, meaning the 8088 got more usage and grew more popular. Due to this popularity, Intel innovated upon the 8088, creating new and improved CPUs that ran based on 32-bit (and eventually 64-bit) instructions, maintaining the 86 suffix on the naming [4].

With a modern strategy of maintaining backwards compatibility and a need to make instructions faster, Intel ditched the numerical naming for more standard branding such as Pentium and Centrino. With the 8086 being the start of a trend of rapid developments and add ones to the x86 architecture, Intel's CPUs holds a presence today like never before. x86 is the basis of most computer architectures today, all due to IBM deciding to run with it for their PC.

### 3 Control Structures

x86 consists of one primary control-flow instruction called jump. There are also a variety of other instructions that expand upon from the jump instruction, such as loop, leave, and call instructions. Although there are multiple control structures in x86, they will not be included as they all utilize the jump instruction.

There are two primary types of jumps in x86, conditional and unconditional. Conditional jumps are executed in a very similar way to that of BASIC's GOTO or C if statements, and unconditional, jumps which will always execute. Unconditional jumps, which are represented in x86 with the instruction JMP, requires one label. A label is a string of text followed by a colon, usually used to represent a new segment of code and which allows the jump instruction to move control of the program from the current location to a new location within the program. Jumps are almost unrestricted in their movement throughout the program as they can traverse to locations earlier or further within the code.

Conditional jumps operate in a similar manner but require a flag and a variety of different jump instructions. The flag for the jump is set through the instruction compare, CMP. This instruction compares two values, either originating from registers or integer constants, and sets a flag value in the EFLAGS register according the result of the CMP instruction [2] from Appendix B. Once the flag has been set from the CMP instruction, there are multiple conditions that can be utilized to determine if a jump should happen. The following major jump instructions include:

JE	Jump if equal
JNE	Jump not equal,
JG	Jump if greater than
JGE	Jump if greater than or equal
JL	Jump less than
JLE	Jump less than equal
JS	Jump if negative
JNS	Jump if not negative

[NOTE] This list does not include all the possible jump instructions but represent instructions widely used, all jump types can be found on the official INTEL x86 User Manual, Volume 1. Sections 5-7 and 7-16 [2].

Depending on which jump is used and the result of the CMP instruction, the program will either fall through, which is equivalent to the condition of an if statement being false, or jump to the corresponding label, equivalent to a true condition. It is important to know that jumps are limited to how far they can jump. Both 64 and 32 bit versions of x86, are limited to a maximum of 32 bit offsets from the current execution line [2] Section 7-14 Header 7.3.8.1.

```
_start:
    ; Unconditional Jump
    JMP     _jmp_end

    ; Conditional Jump
    MOV     rax, 0xa
    CMP     rax, 0xa    ; Compare if 10 == 10
    JEQ     _jeq_end    ; Jump if result from compare is equal

_jmp_end:
    ...
```

```
_jeq_end :
    ...
```

Although the following instructions are not necessarily unique control-flow statements, as they are shorthand for the jump instruction, they are never the less useful. The first instruction is LOOP, which does exactly what it seems, iterates while a condition is true. The loop utilizes the RCX register as a counter for the loop, equivalent to *i* in a for loop [2] Sections 7-16 and 7-17 Header 7.3.8.2. Additionally it utilizes a label in order to know where to jump after the loop body has finished executing. To determine when to exit the loop the following instructions can be used:

LOOPE	Loop while equal
LOOPNE	Loop while not equal
LOOPZ	Loop while zero
LOOPNZ	Jump while not zero

Depending on which LOOP instruction is used from the list above, a compare will happen between the RCX register, which defaults to decrements after each loop iteration. If the selected loop instruction determines that the loop has finished executing, the loop will fall through and continue to the line immediately following the LOOP instruction.

The same idea is applied to the string repeat instruction, or REP. However, instead of using the RCX register, the RSI and RDI registers are decremented or incremented to point to the next byte or word in the string [2] Section 7-18 Header 7.3.9.1]. To determine whether or not the registers are incremented or decremented, the EFLAGS will be set to 0 for incremented and 1 for decremented [2] Section 7-19 Header 7.3.9.2.

```
_start :
    MOV     rcx , 10      ; Run loop 10 times

_loop_head :
    ; Loop Body

    ...

    LOOP   _loop_head    ; Will loop as long as rcx >= 0
```

## 4 Data Types

x86 provides a multitude of fundamental data types upon which can be built on to develop much more sophisticated data structures. Both 32-bit and 64-bit x86 versions contain the ability to use bytes, words, and doublewords. A byte is allocated 8 bits, a word is allocated 16 bits (2 bytes), and a doubleword is allocated 32 bits (4 bytes). Once the IA-32 architecture was introduced, quadwords and double quadwords were added to the list of fundamental data types, being able to access 64 bits (8 bytes) and 128 bits (16 bytes) respectively [2] Section 4.1. In memory, words, double words, and quadwords do not need to be aligned on natural boundaries, but can instead be placed wherever, but to be efficient, one should aim to place these data types within even address, as well as addresses divisible by 4 and 8.

x86 serves as a weakly typed language, giving the programmer freedom to re assign registers and variables as needed. When assigning, x86 often requires a prefix before the register denoting what size is to be moved. For example:

```
section .data :
    foo dq 100

section .text :
    MOV al, byte foo
```

will take a quadword and move a byte of the quadword into the 8 bit register *al*. Note that this is taken from the least significant end of the variable.

Within each of the fundamental data types, x86 supports additional interpretations of these types, with the most common being numeric, pointer, bit field, and string data types. Numeric data types are integers as well as floating point numbers. Integers are represented as signed or unsigned values, with unsigned ranging from 0 to  $2^n - 1$  for *n* bits, and signed ranging from  $-2^{n-1}$  to  $2^{n-1} - 1$  for *n* bits [2] Section 4.2 Header 4.2.1.1. Floating point numbers are represented in single precision, double precision, and double-extended precision.

String data types in x86 make special use of bytes to function. By allocating each ASCII character a single byte (requires more if using UTF-8 formatting), strings are able to be represented as an array of bytes, words, or double words [2] Section 4.5. Statically, they can be allocated at compile time through the .data section

```
section .data:
    hello    db    "Hello!"
```

If one wishes to allocate them dynamically, this can be achieved through an allocated uninitialized byte array in the .bss section. This can then be filled in throughout the program.

With these basic fundamental data types, composite types can be built. It is up to the programmer to declare more complicated data structures such as arrays and structs. Arrays in x86 can be declared much exactly like a string (a string simply being a byte array), and can be any fundamental data type. Using such arrays, it is possible to create 2D arrays, and build from that to develop other data structures.

## 5 Subprograms

x86 is primarily composed of procedures, or functions, and other x86 assembly files, similar to that of C files. Procedures, which are usually defined by labels, can have two types of scopes, which may be dependent with different assemblers, local and global. Additional x86 files can be created and linked to the main assembly file in order help modularize code into easy to maintain chunks, which is important especially when dealing with assembly code.

Procedures, sometimes referred to as subroutines, are similar to that of functions in the C programming language. Procedures contain a unique label, accepts an arbitrary number of parameters, contains a code body, and returns a value, the return value is stored within the RAX register by convention. It is important to understand that a procedure is just a label followed by a body of code. Similar to that of C and Java, the NASM assembler has support for local procedures. Local procedures, which follow the form of dot label, .label\_name, allow for multiple same name procedures within different global labels, or labels not beginning with a dot. If there exists multiple same name local procedures, each duplicate name procedure are associated with the previous non-local label [6]. These local procedures can however be called outside the non-locally associated label anywhere in the program by specifying the name of the local procedures non-local label followed by the local procedure name, non\_local\_label.local\_label. In order to execute a procedure, the instruction CALL is used. This instruction is followed by the procedure that is desired, such as CALL calc\_fib\_seq, which calls the procedure that calculates the Fibonacci sequence. CALL utilizes the standard calling convention registers and the result of the procedure is stored in the RAX register. This register may or may not contain a result, depending on the procedure, but the result is guaranteed to be stored in this register [2].

### Standard Procedure

```
MOV     rdi , 10           ; Procedure parameter
CALL    procedure

procedure:
    ADD     rdi , rdi
    MOV     rax , rdi       ; Return value is stored in RAX register
    RET                     ; Jumps to line immediately following the CALLER
```

### Local Procedures

```
local1:
.local_procedure:
    ...
    RET

local2:
.local_procedure:
    ...
    RET
```

[NOTE] The above procedure shows two identical procedures the the procedure directly following the local1 label can be called by using the expanded form local1.local\_procedure while the procedure following the local2 label can be called using local2.local\_procedure.

Additional assembly files can be linked as an additional way to modularize. The way in which additional files are linked together differs based on which assembler is being utilized. Some assemblers require .inc files, similar to that of C header files, or allow assembly files to be included into the programming that will be calling the file by using `%include "<assembly file>"`. The layout of additional files is very similar to that of most languages that support multi-file programs. Each additional assembly file, which may or may not need a corresponding .inc file, only have a few restrictions to what may be contained within them. The most important restriction is that each assembly file must not contain the global label `_start`, similar to C's main function, and duplicate label names across files should be avoided, but may be used depending on a variety of cases such as local label names [5]. The keyword `EXTERN`, can also be used to call libraries outside the scope of the program, such as C libraries that preexist on the local machine.

```
EXTERN printf
#include "additional_file.asm"

section .text
...
```

## 6 Summary

x86 has lead the way in CISC architecture due to its simplistic nature and how performance power. With the ability to power 16-bit, 32-bit, and 64-bit machines and the ability of backwards compatibility with smaller bit machines, x86 is versatile and verbose. As compilers ranging from Microsoft to open source such as GNU utilize x86 assembly, the language continues to be expanded upon added more complex and efficient instructions. Although assembly can be found almost anywhere, the language is not often utilized by developers but rather compilers. As a result of this, x86 evolved into different syntax's, Intel for more DOS and Windows machine and ATT for Unix, but Intel is a preferred syntax when developers require to read or develop. To this day, assembly is replaced for more convenience and time efficient languages that allow for easier development while maintaining most of the efficiency that x86 provides.

## References

- [1] Benj Edwards. Birth of a standard: The intel 8086 microprocessor turns 40 today. [www.pcworld.com/article/146957/article.html](http://www.pcworld.com/article/146957/article.html), 2018. Online.
- [2] Intel Corporation. *Intel © 64 and IA-32 Architectures Software Developer's Manual*. Intel, 2019.
- [3] Emily Blem, Jaikrishnan Menon, and Karthikeyan Sankaralignam. Cpu comparison: x86 vs arm - will intel i9 9900k stay atop? <https://research.cs.wisc.edu/vertical/papers/2013/hpca13-isa-power-struggles.pdf>, 2013. 19th IEEE International Symposium on High Performance Computer Architecture.
- [4] Intel x86 Wiki. Intel x86 - computer history wiki, computer history wiki. [gunkies.org/wiki/Intelx86](http://gunkies.org/wiki/Intelx86), 2018. Online.
- [5] NASM. The netwide assembler: Nasm. <https://www.nasm.us/xdoc/2.09.04/html/nasmdoc1.html>. Online.
- [6] Tortall Networks. The nasm language local labels. <https://www.tortall.net/projects/yasm/manual/html/nasm-local-label.html>. Online.